

Write a program to predict the eligibility of a customer for loan disbursement.

Process

1. Getting the system ready and loading the data
2. Understanding the data
3. Model Building: Part 1
4. Logistic Regression using stratified k-folds cross-validation

Important Libraries

Python

Pandas

matplotlib

seaborn

sklearn

In []:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

In []:

```
from google.colab import files
uploaded=files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session.
Please rerun this cell to enable.

Saving test.csv to test.csv

Saving train.csv to train.csv

In []:

```
train=pd.read_csv("train.csv")
train.head()
```

Out[3]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	

In []:

```
test=pd.read_csv("test.csv")
test.head()
```

Out[4]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome
0	LP001015	Male	Yes	0	Graduate	No	5720	
1	LP001022	Male	Yes	1	Graduate	No	3076	
2	LP001031	Male	Yes	2	Graduate	No	5000	
3	LP001035	Male	Yes	2	Graduate	No	2340	
4	LP001051	Male	No	0	Not Graduate	No	3276	

In []:

```
train.columns
```

Out[5]:

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
      'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
      'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

In []:

```
print(train.shape, test.shape)
```


(614, 13) (367, 12)

In []:

```
train['Loan_Status'].value_counts().plot.bar()
```

Out[7]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7911570990>



Categorical features: These features have categories (Gender, Married ,Self_Employed, Credit_History, Loan_Status)

Ordinal features: Variables in categorical features having some order involved (Dependents, Education, Property_Area)


Numerical features: These features have numerical values (ApplicantIncome, CoapplicantIncome, LoanAmount, Loan_Amount_Term)

In []:

```
train['Gender'].value_counts(normalize=True).plot.bar(title='Gender')
```

Out[8]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7910db6210>




In []:

```
train['Married'].value_counts(normalize=True).plot.bar(title='Married')
```

Out[9]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7911479f90>




In []:

```
train['Self_Employed'].value_counts(normalize=True).plot.bar(title='Self_employed')
```

Out[10]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7910545a50>




In []:

```
train['Credit_History'].value_counts(normalize=True).plot.bar(title='Credit_History')
```

Out[11]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7910525b90>




In []:

```
train['Dependents'].value_counts(normalize=True).plot.bar(title='Dependents')
```

Out[12]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f79103fdc10>




In []:

```
train['Education'].value_counts(normalize=True).plot.bar(title='Education')
```

Out[13]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7910354190>




In []:

```
train['Property_Area'].value_counts(normalize=True).plot.bar(title='Property_Area')
```

Out[14]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f791034ab50>




In []:

```
train['ApplicantIncome'].value_counts(normalize=True).plot.bar(title='ApplicantIncome')
```

Out[15]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f79102cc950>




In []:

```
train['CoapplicantIncome'].plot.box()
```

Out[16]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f7910252210>

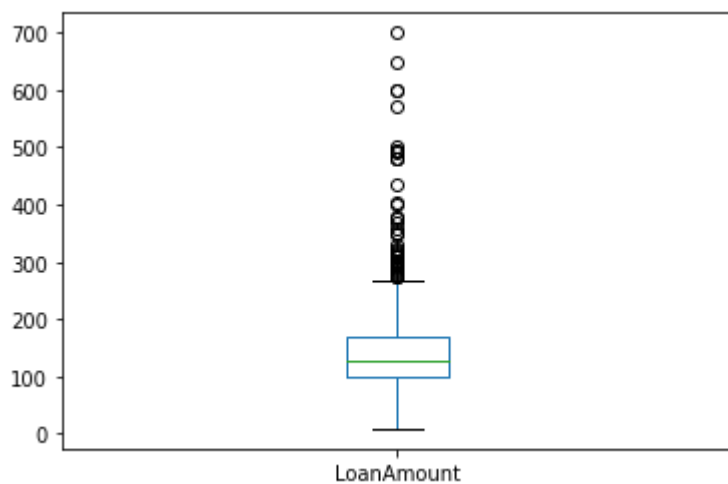


In []:

```
train['LoanAmount'].plot.box()
```

Out[17]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f790f881490>




In []:

```
Gender=pd.crosstab(train['Gender'],train['Loan_Status'])
Gender.div(Gender.sum(1).astype(float),axis=0).plot(kind='bar',stacked=True,figsize=(4,4))
```

Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x7fe24026ce10>



In []:

```
matrix=train.corr()  
f,ax=plt.subplots(figsize=(9,6))  
sns.heatmap(matrix,vmax=.8,square=True,cmap="BuPu",annot=True)
```

Out[18]:

<matplotlib.axes._subplots.AxesSubplot at 0x7f790f655610>



In []:

```
train.isnull().sum()
```

Out[20]:

```
Loan_ID          0
Gender           13
Married          3
Dependents       15
Education        0
Self_Employed   32
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       22
Loan_Amount_Term 14
Credit_History  50
Property_Area    0
Loan_Status      0
dtype: int64
```

In []:

```
train['Gender'].fillna(train['Gender'].mode()[0],inplace=True)
print(train['Gender'])
```

```
0      Male
1      Male
2      Male
3      Male
4      Male
...
609    Female
610     Male
611     Male
612     Male
613    Female
Name: Gender, Length: 614, dtype: object
```

In []:

```
train['Married'].fillna(train['Married'].mode()[0],inplace=True)
print(train['Married'])
```

```
0      No
1      Yes
2      Yes
3      Yes
4      No
...
609    No
610    Yes
611    Yes
612    Yes
613    No
Name: Married, Length: 614, dtype: object
```

In []:

```
train['Dependents'].fillna(train['Dependents'].mode()[0],inplace=True)
print(train['Dependents'])
```

```
0      0
1      1
2      0
3      0
4      0
...
609    0
610   3+
611    1
612    2
613    0
Name: Dependents, Length: 614, dtype: object
```

In []:

```
train['Self_Employed'].fillna(train['Self_Employed'].mode()[0],inplace=True)
print(train['Self_Employed'])
```

```
0      No
1      No
2     Yes
3      No
4      No
...
609    No
610    No
611    No
612    No
613   Yes
Name: Self_Employed, Length: 614, dtype: object
```

In []:

```
train['LoanAmount'].fillna(train['LoanAmount'].mode()[0],inplace=True)
print(train['LoanAmount'])
```

```
0     120.0
1     128.0
2      66.0
3     120.0
4     141.0
...
609    71.0
610    40.0
611   253.0
612   187.0
613   133.0
Name: LoanAmount, Length: 614, dtype: float64
```

In []:

```
train['Loan_Amount_Term'].fillna(train['Loan_Amount_Term'].mean(),inplace=True)
print(train['LoanAmount'])
```

```
0      120.0
1      128.0
2       66.0
3      120.0
4      141.0
...
609     71.0
610     40.0
611    253.0
612    187.0
613    133.0
Name: LoanAmount, Length: 614, dtype: float64
```

In []:

```
train['Credit_History'].fillna(train['Credit_History'].mean(),inplace=True)
print(train['Credit_History'])
```

```
0      1.0
1      1.0
2      1.0
3      1.0
4      1.0
...
609     1.0
610     1.0
611     1.0
612     1.0
613     0.0
Name: Credit_History, Length: 614, dtype: float64
```

In []:

```
train.isnull().sum()
```

Out[28]:

```
Loan_ID      0
Gender        0
Married       0
Dependents    0
Education     0
Self_Employed 0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
Loan_Status   0
dtype: int64
```

In []:

```
X=train.drop('Loan_Status',1)
#X=train.drop('Loan_ID')
X=X.drop('Loan_ID',axis=1)
Y=train['Loan_Status']

X=pd.get_dummies(X)

print(X)
print(Y)
```

	ApplicantIncome	...	Property_Area_Urban
0	5849	...	1
1	4583	...	0
2	3000	...	1
3	2583	...	1
4	6000	...	1
..
609	2900	...	0
610	4106	...	0
611	8072	...	1
612	7583	...	1
613	4583	...	0

[614 rows x 20 columns]

0	Y
1	N
2	Y
3	Y
4	Y
..	
609	Y
610	Y
611	Y
612	Y
613	N

Name: Loan_Status, Length: 614, dtype: object

In []:

```
from sklearn.linear_model import LogisticRegression
model=LogisticRegression()
clf=model.fit(X,Y)
df=pd.read_csv('test.csv')
print(df)
```

	Loan_ID	Gender	Married	...	Loan_Amount_Term	Credit_History	Property_A
rea							
0	LP001015	Male	Yes	...	360.0	1.0	Ur
ban							
1	LP001022	Male	Yes	...	360.0	1.0	Ur
ban							
2	LP001031	Male	Yes	...	360.0	1.0	Ur
ban							
3	LP001035	Male	Yes	...	360.0	NaN	Ur
ban							
4	LP001051	Male	No	...	360.0	1.0	Ur
ban							
..	
...							
362	LP002971	Male	Yes	...	360.0	1.0	Ur
ban							
363	LP002975	Male	Yes	...	360.0	1.0	Ur
ban							
364	LP002980	Male	No	...	360.0	NaN	Semiur
ban							
365	LP002986	Male	Yes	...	360.0	1.0	Ru
ral							
366	LP002989	Male	No	...	180.0	1.0	Ru
ral							

[367 rows x 12 columns]

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:94
0: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

In []:

```
df.isnull().sum()
```

Out[31]:

```
Loan_ID          0
Gender           11
Married          0
Dependents       10
Education         0
Self_Employed    23
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       5
Loan_Amount_Term 6
Credit_History  29
Property_Area    0
dtype: int64
```

In []:

```
df['Gender'].fillna(df['Gender'].mode()[0],inplace=True)
df['Dependents'].fillna(df['Dependents'].mode()[0],inplace=True)
df['Self_Employed'].fillna(df['Self_Employed'].mode()[0],inplace=True)
df['LoanAmount'].fillna(df['LoanAmount'].mean(),inplace=True)
#df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].mean(),inplace=True)
df['Credit_History'].fillna(df['Credit_History'].mean(),inplace=True)
```

In []:

```
df.isnull().sum()
df['Loan_Amount_Term'].fillna(df['Loan_Amount_Term'].mean(),inplace=True)
```

In []:

```
df.isnull().sum()
```

Out[34]:

```
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education         0
Self_Employed    0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History   0
Property_Area    0
dtype: int64
```


In []:

```
df=df.drop('Loan_ID',axis=1)
df=pd.get_dummies(df)
print(X)
print(df)
```

	ApplicantIncome	...	Property_Area_Urban
0	5849	...	1
1	4583	...	0
2	3000	...	1
3	2583	...	1
4	6000	...	1
..
609	2900	...	0
610	4106	...	0
611	8072	...	1
612	7583	...	1
613	4583	...	0

[614 rows x 20 columns]

	ApplicantIncome	...	Property_Area_Urban
0	5720	...	1
1	3076	...	1
2	5000	...	1
3	2340	...	1
4	3276	...	1
..
362	4009	...	1
363	4158	...	1
364	3250	...	0
365	5000	...	0
366	9200	...	0

[367 rows x 20 columns]

In []:

```
clf.predict(df)
```

Out[36]:

[illegible]